

LOGISTICAL REGRESSION

In [1]:

```
#Aim: To perform Logistical Regression
#Exp no:10
#Name:Khushi Chandrashekhar Satpute
#Sec:B
#Roll no:43
#Sub:ET-1
#Date:11/10/2024
```

Importing Libraries

In [4]:

```
import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns
from sklearn.model_selection import train_test_split
import warnings
warnings.filterwarnings('ignore')
```

In [5]:

```
import os
```

In [6]:

```
os.getcwd()
```

Out[6]:

```
'C:\\Users\\asus'
```

In [7]:

```
os.chdir("C:\\Users\\asus\\Desktop")
```

In [8]:

```
df=pd.read_csv("framingham.csv")
```

In [9]:

```
#The "Framingham" heart disease dataset includes over 4,240 records, 15 attributes.
#The goal of the dataset is to predict whether the patient has 10-year risk of future (C
```

In [10]:

```
df.head()
```

Out[10]:

	male	age	education	currentSmoker	cigsPerDay	BPMeds	prevalentStroke	prevalentHyp	diabetes
0	1	39	4.0	0	0.0	0.0	0	0	0
1	0	46	2.0	0	0.0	0.0	0	0	0
2	1	48	1.0	1	20.0	0.0	0	0	0
3	0	61	3.0	1	30.0	0.0	0	1	0
4	0	46	3.0	1	23.0	0.0	0	0	0

In [11]:

```
df.describe()
```

Out[11]:

	male	age	education	currentSmoker	cigsPerDay	BPMeds	prevalentStroke
count	4238.000000	4238.000000	4133.000000	4238.000000	4209.000000	4185.000000	4238.000000
mean	0.429212	49.584946	1.978950	0.494101	9.003089	0.029630	0.00589
std	0.495022	8.572160	1.019791	0.500024	11.920094	0.169584	0.07658
min	0.000000	32.000000	1.000000	0.000000	0.000000	0.000000	0.00000
25%	0.000000	42.000000	1.000000	0.000000	0.000000	0.000000	0.00000
50%	0.000000	49.000000	2.000000	0.000000	0.000000	0.000000	0.00000
75%	1.000000	56.000000	3.000000	1.000000	20.000000	0.000000	0.00000
max	1.000000	70.000000	4.000000	1.000000	70.000000	1.000000	1.00000

In [12]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4238 entries, 0 to 4237
Data columns (total 16 columns):
#   Column                Non-Null Count  Dtype
---  -
0   male                  4238 non-null  int64
1   age                   4238 non-null  int64
2   education             4133 non-null  float64
3   currentSmoker        4238 non-null  int64
4   cigsPerDay            4209 non-null  float64
5   BPMeds                4185 non-null  float64
6   prevalentStroke       4238 non-null  int64
7   prevalentHyp          4238 non-null  int64
8   diabetes              4238 non-null  int64
9   totChol              4188 non-null  float64
10  sysBP                 4238 non-null  float64
11  diaBP                 4238 non-null  float64
12  BMI                   4219 non-null  float64
13  heartRate             4237 non-null  float64
14  glucose               3850 non-null  float64
15  TenYearCHD            4238 non-null  int64
dtypes: float64(9), int64(7)
memory usage: 529.9 KB
```

In [13]:

```
df.isna().sum()
```

Out[13]:

```
male          0
age           0
education     105
currentSmoker 0
cigsPerDay    29
BPMeds        53
prevalentStroke 0
prevalentHyp  0
diabetes      0
totChol       50
sysBP         0
diaBP         0
BMI           19
heartRate     1
glucose       388
TenYearCHD    0
dtype: int64
```

In [14]:

```
#Since, only a few rows have null values in them, we are only removing those rows from the
df = df.dropna(subset=['heartRate', 'BMI', 'cigsPerDay', 'totChol', 'BPMeds'])
```

In [15]:

```
df
```

Out[15]:

	male	age	education	currentSmoker	cigsPerDay	BPMeds	prevalentStroke	prevalentHyp	diabetes
0	1	39	4.0	0	0.0	0.0	0	0	
1	0	46	2.0	0	0.0	0.0	0	0	
2	1	48	1.0	1	20.0	0.0	0	0	
3	0	61	3.0	1	30.0	0.0	0	1	
4	0	46	3.0	1	23.0	0.0	0	0	
...
4233	1	50	1.0	1	1.0	0.0	0	1	
4234	1	51	3.0	1	43.0	0.0	0	0	
4235	0	48	2.0	1	20.0	NaN	0	0	
4236	0	44	1.0	1	15.0	0.0	0	0	
4237	0	52	2.0	0	0.0	0.0	0	0	

4238 rows × 10 columns

Missing Value Treatment

Since, 'glucose' and 'education' columns had a significant amount of null values, so we replaced them with the mean of values for their respective columns

In [30]:

```
df['glucose'].fillna(value = df['glucose'].mean(),inplace=True)
```

In [31]:

```
df['education'].fillna(value = df['education'].mean(),inplace=True)
```

In [32]:

```
df['heartRate'].fillna(value = df['heartRate'].mean(),inplace=True)
```

In [36]:

```
df['BMI'].fillna(value = df['BMI'].mean(),inplace=True)
```

In [38]:

```
df['cigsPerDay'].fillna(value = df['cigsPerDay'].mean(),inplace=True)
```

In [40]:

```
df['totChol'].fillna(value = df['totChol'].mean(),inplace=True)
```

In [42]:

```
df['BPMeds'].fillna(value = df['BPMeds'].mean(),inplace=True)
```

In [44]:

```
df.isna().sum()
```

Out[44]:

male	0
age	0
education	0
currentSmoker	0
cigsPerDay	0
BPMeds	0
prevalentStroke	0
prevalentHyp	0
diabetes	0
totChol	0
sysBP	0
diaBP	0
BMI	0
heartRate	0
glucose	0
TenYearCHD	0

dtype: int64

In [46]:

```
#Splitting the dependent and independent variables.  
x = df.drop("TenYearCHD",axis=1)  
y = df['TenYearCHD']
```

In [48]:

```
x #checking the features
```

Out[48]:

	male	age	education	currentSmoker	cigsPerDay	BPMeds	prevalentStroke	prevalentHyp	diabe
0	1	39	4.0	0	0.0	0.00000	0	0	
1	0	46	2.0	0	0.0	0.00000	0	0	
2	1	48	1.0	1	20.0	0.00000	0	0	
3	0	61	3.0	1	30.0	0.00000	0	1	
4	0	46	3.0	1	23.0	0.00000	0	0	
...
4233	1	50	1.0	1	1.0	0.00000	0	1	
4234	1	51	3.0	1	43.0	0.00000	0	0	
4235	0	48	2.0	1	20.0	0.02963	0	0	
4236	0	44	1.0	1	15.0	0.00000	0	0	
4237	0	52	2.0	0	0.0	0.00000	0	0	

4238 rows × 15 columns

Train Test Split

In [51]:

```
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.2,random_state=42)
```

In [53]:

```
y_train
```

Out[53]:

```
3252    0
3946    0
1261    0
2536    0
4089    0
..
3444    0
466     0
3092    0
3772    0
860     0
```

Name: TenYearCHD, Length: 3390, dtype: int64

Logistic Regression Algorithm

In [56]:

```
from sklearn.linear_model import LogisticRegression
model = LogisticRegression().fit(x_train,y_train)
model.score(x_train, y_train)
```

Out[56]:

```
0.848377581120944
```

In []: